

Part A: Convolutional Neural Networks (Image Classification)

submitted by Manjusha Deore

Introduction

This report presents a comprehensive investigation into the training, evaluation, and robustness of Convolutional Neural Networks (CNNs) for image classification. It compares custom architectures trained from scratch against transfer learning techniques using a pre-trained VGG16 network, systematically evaluating performance under clean conditions and additive Gaussian noise disturbances on the CIFAR-100 dataset.

Dataset Description

The experiments utilize the **CIFAR-100** dataset, consisting of 100 fine-grained object categories. Each sample is a 32×32 RGB image. The dataset is pre-partitioned into 50,000 training images and 10,000 test images. Training a model from scratch on 100 classes with this low spatial resolution presents significant risks of overfitting, requiring robust architectural design.

1 Question A1: Baseline CNN Architecture and Training

1.1 Architecture Design Choice

For CIFAR-100, we require a model with capacity to capture complex features but small enough to train reasonably fast.

- **Structure:** A VGG-style architecture with 3 convolutional blocks.
- **Components:** Each block uses 3×3 Convolutional layers, Batch Normalization (to stabilize and speed up training), ReLU activations, and 2×2 Max Pooling.
- **Classifier:** Fully connected layers with Dropout to prevent overfitting.

Architectural Choices: 3×3 kernels were selected as they are standard for capturing local spatial patterns without causing a massive explosion in parameters. Batch Normalization prevents vanishing/exploding gradients.

1.2 Data Processing & Optimization

Data Processing: The pipeline explicitly incorporates `RandomCrop` and `RandomHorizontalFlip`. Without these spatial augmentations, training from scratch results in severe overfitting.

Optimization: The Adam optimizer was combined with `weight_decay` (L2 Regularization) to navigate the complex loss landscape while penalizing large weights to combat overfitting.

1.3 Convergence Behavior & Results

The training loss steadily decreases, but the validation loss plateaus after 15-20 epochs, indicating the limits in generalizing from scratch.




Figure 1: Training and Validation Loss and Accuracy over 40 Epochs.

Performance Metric	Result
Final Training Loss	1.0511
Final Training Accuracy	69.43%
Final Validation Loss	1.2439
Final Validation Accuracy	65.35%
Best Validation Accuracy (Hit at Epoch 38)	65.42%

Table 1: Detailed Training Performance Report (A1)

Overall Performance & Generalization: The model achieved a Validation Accuracy of 64.21% (averaging runs). The Generalization Gap was crushed to just 4.9% by increasing weight decay to $1e-3$. The model sacrificed its ability to "memorize" training data to learn generalizable features. A $\sim 64\%$ accuracy on CIFAR-100 from scratch with only a 5% gap is an exceptionally healthy baseline.

2 Question A2: Robustness to Additive Gaussian Noise

2.1 Methodology and Design Decisions

Math of the Noise Pipeline: Since the problem gave a variance of $\sigma^2 = 0.05$, the standard deviation was calculated dynamically using $\sigma = \sqrt{0.05} \approx 0.2236$.

Clamping Justification: `torch.clamp(noisy_tensor, 0.0, 1.0)` was necessary because raw additive noise could generate physically impossible pixel values (< 0 or > 1). Clamping ensures test images remain valid representations.

Documentation Strategy: Noise parameters were encapsulated into a JSON file (`noise_schedule.json`) to ensure reproducibility for subsequent experiments.

2.2 Expected Drop in Performance & Results

It was expected that accuracy would plummet. A standard CNN trained purely on clean data relies heavily on high-frequency spatial patterns, which Gaussian noise destroys.

Test Condition	Cross-Entropy Loss	Accuracy (%)
Clean Data (Baseline A1)	1.2439	65.35%
Noisy Data (Variance=0.05)	11.2236	1.40%

Table 2: Detailed Performance Table

Metric	Value	Interpretation
Accuracy Drop (Absolute)	-63.95%	Total percentage points lost
Accuracy Drop (Relative)	-97.86%	Percentage of original accuracy lost
Loss Increase (Absolute)	+9.9797	Increase in model uncertainty

Table 3: Degradation Analysis Table

3 Question A3: Improved Training Strategy

3.1 The Strategy: Probabilistic Fine-Tuning

To satisfy the assignment's constraint of using a "limited number" of perturbed samples while keeping the distribution "close to the clean dataset," **Probabilistic Noise Injection via Fine-Tuning** was used.

The Method (RandomApply): PyTorch's `RandomApply` injected the exact Gaussian noise with a 20% probability ($p = 0.2$). Thus, 80% of training images remained perfectly clean. The training distribution is mathematically anchored to the clean dataset, but the model is forced to learn noise-invariant features.

Fine-Tuning vs. Scratch: The converged A1 weights were loaded and fine-tuned for 10-15 epochs with a lower learning rate (1×10^{-4}). This targeted regularization phase allows established feature-extractors to "adapt" stably.

Model Version	Clean Acc.	Noisy Acc. (Var=0.05)	Degradation Gap
Baseline (A1)	65.35%	1.37%	-63.98%
Robust Model (A3)	64.47%	28.82%	-35.65%

Table 4: Robustness Recovery

4 Question A4: Transfer Learning with Pre-trained VGG16

4.1 Methodology

Classifier Dimension Adjustment: VGG expects 224×224 images. CIFAR-100 32×32 images pass through pooling layers and reduce to $1 \times 1 \times 512$. The new MLP classifier is built to accept exactly 512 inputs.

Computational Efficiency: Only `model.classifier.parameters()` was passed to the optimizer. The convolutional base was locked (`requires_grad = False`). This mathematically shorter backward pass saved massive GPU memory and time.

Data Efficiency: The pre-trained VGG model already learned universal visual features from ImageNet, acting as a powerful "lens" for the tiny MLP.

Domain Shift Caveat: Accuracy jumps significantly, but doesn't reach 90%+ because CIFAR-100 (32×32 blurry images) has a "domain shift" compared to ImageNet (high-resolution photos).

Architecture Setup	Training Approach	Test Acc (%)
Baseline Custom CNN (A1)	Scratch (40 Epochs)	64.21%
Pre-trained VGG16 + MLP (A4)	Strict Transfer Learning (Frozen VGG Base)	69.22%

Table 5: Performance Gain via Transfer Learning

5 Question A5: Robustness of Transfer Learning

The transfer learning model was evaluated against the identical noise disturbance.

Model	Clean Accuracy	Noisy Accuracy ($\sigma^2 = 0.05$)	Performance Drop
CNN (Scratch)	64.21%	12.50%	Lost 80.5% capacity
VGG16 (Transfer)	69.22%	3.21%	Lost 95.4% capacity

Table 6: Robustness Comparison Report




Figure 2: Comparative analysis of model robustness.

6 Question A6: Comprehensive Discussion and Conclusion

6.1 Detailed Observations and Challenges

A. Training from Scratch (A1): Training on CIFAR-100 is inherently difficult due to tiny 32×32 resolution. The primary challenge was overfitting. We implemented severe regularization: Dropout, Batch Normalization, heavy Data Augmentation, and Cosine Annealing to navigate the loss landscape.

B. Noise Injection (A2): The baseline CNN suffered a catastrophic degradation (accuracy plummeted to $\sim 1.4\%$). Injecting raw noise pushed pixel values outside the visual spectrum, solved by `torch.clamp`. Theoretically, the CNN over-relied on high-frequency spatial textures which the noise obliterated.

C. Mitigating Noise (A3): The challenge was "catastrophic forgetting" of clean data. We solved this via Probabilistic Noise Injection ($p = 0.20$). The noise acted as a regularizer, forcing noise-invariant feature learning while anchored to clean data.

D. Transfer Learning (A4 & A5): The major challenge was Resolution Blindness. Upscaling to 224×224 caused CUDA Out-of-Memory (OOM) crashes. We bypassed the upscale entirely using `nn.AdaptiveAvgPool2d((1, 1))` to cleanly squash spatial dimensions into a $1 \times 1 \times 512$ feature vector, preventing crashes and keeping training blazing fast.

6.2 Comparison of Approaches

Classification Performance: Transfer Learning using VGG16 achieved vastly superior performance. The pre-trained early layers acted as perfect edge/color detectors, meaning the small MLP only had to learn CIFAR class mappings. The From-Scratch CNN was bottlenecked by limited data.

Robustness to Noise: The From-Scratch CNN exhibited severe brittleness without explicit training (A3). Transfer Learning exhibited unexpected degradation here due to the extreme domain shift between ImageNet and noisy 32×32 static, showing that even robust extractors require targeted fine-tuning for adversarial distributions.

Computational Efficiency: The From-Scratch CNN was highly inefficient, taking considerable GPU time for millions of parameters over 40 epochs. Transfer Learning was highly efficient; freezing the base skipped the heaviest backward passes, meaning a full epoch took mere seconds.

6.3 Conclusion: The Best Trade-Off

Transfer Learning using a frozen pre-trained base provides the absolute best trade-off between computational efficiency and clean classification performance. Training a CNN from scratch is a massive expenditure of resources for mediocre results. By locking pre-trained weights, training time is reduced to a fraction of the baseline, yielding superior clean accuracy, making it the clear winner for computationally constrained real-world deployment.

Part B-Comparative Analysis of Recurrent Neural Architectures for Sentiment Classification

Manjusha Deore

1 Introduction

The objective of this assignment is to study and compare the effectiveness of different recurrent neural network architectures, namely vanilla recurrent neural networks (RNNs), long short-term memory networks (LSTMs), and attention-augmented variants, for binary sentiment classification. This report details the implementation, challenges, and performance metrics of these models using various word embedding strategies.

2 Dataset Description

Large Movie Review Dataset: This is a dataset for binary sentiment classification containing substantially more data than previous benchmark datasets. We provide a set of 25,000 highly polar movie reviews for training, and 25,000 for testing. There is additional unlabeled data for use as well. Raw text and already processed bag of words formats are provided. For the purpose of these experiments, we utilize the version 1.0 of this dataset, specifically targeting the labeled polarized reviews to train models to output a single sentiment label.

3 Question B1: Preprocessing

The preprocessing phase is critical to reduce vocabulary redundancy and prepare sequences for recurrent processing.

- **Normalization:** Standard techniques including lower-casing and punctuation stripping were applied.
- **Tokenization:** Word-to-integer mapping was performed with a vocabulary limit.
- **Sequence Handling:** Optimized for RNNs using **Pre-Padding** and **Pre-Truncating**.

Execution Log Details:

- Loading data from *imdb.npz*.
- Training/Testing samples: 25,000 each.
- Vocabulary Size: 10,000.
- Fixed Sequence Length: 150.
- Sample Tokenization: [12, 16, 43, 530, 38, 76, 15, 13, 1247, 4]

4 Question B2: Vocabulary and Embeddings

4.1 Vocabulary & Mapping Procedure

The vocabulary was constructed by analyzing the IMDB dataset and selecting the top 10,000 most frequent tokens. Each token was mapped to a unique integer index based on frequency rank. We reserved index 0 for padding and index 1 for the start-of-sequence token.

4.2 Embedding Conversion Methods

- **One-Hot (Trainable):** Initially randomized vectors learned during backpropagation. This serves as a dense approximation of one-hot encoding, as modern deep learning avoids sparse $10,000 \times 10,000$ matrices for computational efficiency.
- **Word2Vec:** Pre-trained vectors via the *gensim* library (Google News corpus), capturing semantic relationships via Skip-gram/CBOW.
- **GloVe:** Stanford pre-trained vectors utilizing global word-word co-occurrence statistics.

Global Parameters: Vocab: 10,000; Embedding Dim: 100; Sequence Length: 100; Hidden Units: 64.

5 Question B3: Vanilla RNN Architecture

5.1 Architecture Justification

The architecture was initially set at $L = 2$ hidden layers but later refined to $L = 1$ with $h = 64$.

- **Depth (L):** Vanilla RNNs are difficult to train beyond 2 layers due to gradient instability. $L = 1$ provides stability and avoids the vanishing gradient trap where lower layer weights stop updating.
- **Dimension (h):** 64 units provide capacity for reviews without leading to immediate overfitting.

5.2 Challenges: Vanishing Gradients

The primary challenge is the **Vanishing Gradient**. During backpropagation, gradients are multiplied by weights at every timestep. For long sequences, the gradient shrinks exponentially, causing the model to "forget" the beginning of the review. This limits the model's ability to handle **Long-Range Dependencies**.

Results (Pre-Padding):

- Epoch 1: Accuracy 0.6624, Loss 0.5929.
- Epoch 2: Accuracy 0.8381, Loss 0.3793.
- **Improved Test Accuracy: 0.8117.**

6 Question B4: Embedding Integration**6.1 Architectural Modifications**

To integrate pre-trained embeddings, we modified the Embedding layer initialization:

- **Trainable (One-Hot Equivalent):** Weights set to random; `trainable=True`.
- **Pre-trained (GloVe/Word2Vec):** Weights set from external corpora; `trainable=False` (frozen). This significantly reduces complexity by not updating the 1M parameters in the embedding layer.

6.2 Performance Metrics and Differences

Table 1: B4 Performance Comparison

Method	Acc	Prec	F1
Trainable	0.8171	0.8163	0.8173
Static_GloVe	0.7220	0.8138	0.6744
Static_W2V	0.7214	0.6724	0.7560

6.3 Observed Differences

- **Cold Start:** Trainable models start with random noise, causing lower Epoch 1 accuracy compared to GloVe’s ”warm start”.
- **Generalization:** GloVe/W2V yield higher recall as they understand rare words from external massive corpora.
- **Overfitting:** Trainable models are more prone to overfitting as they tune 1M additional parameters.

7 Question B5: RNN vs. LSTM Comparison

We directly compared the Vanilla RNN with LSTM using $h = 64$ and 5 epochs.

- **Training Stability:** LSTMs utilize a ”Constant Error Carousel” to stabilize gradients, leading to a much lower standard deviation in validation accuracy (0.0055 vs 0.0388).
- **Gating Mechanism:** The LSTM’s Forget Gate allows the network to ”reset” memory, while the Cell State allows gradients to flow linearly, preventing vanishing.

Table 2: B5 Comparison Report

Metric	Vanilla RNN	LSTM
Acc	0.8002	0.8420
F1	0.8078	0.8423
Conv. Speed	23.10s	16.02s
Stability	0.0388	0.0055

8 Question B6: Attention Mechanism**8.1 Feasibility and Justification**

An attention mechanism is highly feasible. It transforms the bottleneck of standard recurrent layers into a dynamic architecture that weights words by importance.

8.2 Implementation Challenges

- **Alignment:** Requires `return_sequences=True`, increasing data volume.
- **Query Selection:** Using the whole sequence as Query and Value adds complexity.
- **Vanishing Weights:** Attention weights can become sparse in very long sequences.

8.3 Computational Overhead

- **Memory:** High (linear with sequence length).
- **Time:** High ($O(N^2)$ attention matrix).

8.4 Rigorous Justification

Attention bypasses the bottleneck where semantic content is compressed into one final hidden state. It computes $C = \sum \alpha_i h_i$, allowing the model to attend to sentiment-heavy tokens regardless of their position.

9 Inference and Verification

Final experiments with Attention-LSTM (Trainable) and LSTM-GloVe showed consistent polar sentiment detection.

Table 3: Sample Verification

Snippet	Prob	Pred	Correct
”...powerful...”	0.9432	Pos	Yes
”...terrible...”	0.3870	Neg	Yes
”...girl...”	0.7991	Pos	No

10 Conclusion

This study confirms that LSTM architectures significantly outperform Vanilla RNNs in stability and long-range dependency handling. While pre-trained embeddings provide

semantic richness, trainable layers offer better adaptability to review-specific slang. Attention mechanisms provide the most robust path to accuracy by removing the recurrent bottleneck.